

A2 - MAD

Julian Wulff

November 2019

Exercise 1

a)

In order to derive the optimal least squares parameter value $\hat{\mathbf{w}}$ of the function:

$$\mathcal{L} = \frac{1}{N} \sum_{n=1}^N \alpha_n (f(\mathbf{x}_n; \mathbf{w}) - t_n)^2 = \frac{1}{N} \sum_{n=1}^N \alpha_n (\mathbf{w}^T \mathbf{x}_n - t_n)^2$$

We will start by rewriting the function in terms of various vectors and matrices, which will be easier to manipulate. We already now how to express

$$\mathcal{L} = \frac{1}{N} \sum_{n=1}^N (\mathbf{w}^T \mathbf{x}_n - t_n)^2 = \frac{1}{N} (\mathbf{X}\mathbf{w} - \mathbf{t})^T (\mathbf{X}\mathbf{w} - \mathbf{t})$$

In order to include α_n in this expression we start by defining a matrix \mathbf{A} as diagonal matrix containing the weights $\alpha_1, \dots, \alpha_n$, in order to multiply these onto our expression. In order to let each individual α_n appear one time each in the final scalar we will multiply \mathbf{A} on to the expression as shown below:

$$\mathcal{L} \frac{1}{N} ((\mathbf{X}\mathbf{w} - \mathbf{t})^T \mathbf{A} (\mathbf{X}\mathbf{w} - \mathbf{t}))$$

Which gives us the desired expression. We will now multiply out the parenthesis in order continue:

$$\mathcal{L} = \frac{1}{N} (\mathbf{X}\mathbf{w} - \mathbf{t})^T \mathbf{A} (\mathbf{X}\mathbf{w} - \mathbf{t}) \quad (1)$$

$$= \frac{1}{N} ((\mathbf{X}\mathbf{w})^T - \mathbf{t}^T) \mathbf{A} (\mathbf{X}\mathbf{w} - \mathbf{t}) \quad (2)$$

$$= \frac{1}{N} ((\mathbf{X}\mathbf{w})^T \mathbf{A} - \mathbf{t}^T \mathbf{A}) (\mathbf{X}\mathbf{w} - \mathbf{t}) \quad (3)$$

$$= \frac{1}{N} (\mathbf{X}\mathbf{w})^T \mathbf{A} \mathbf{X}\mathbf{w} - \frac{1}{N} \mathbf{t}^T \mathbf{A} \mathbf{X}\mathbf{w} - \frac{1}{N} (\mathbf{X}\mathbf{w})^T \mathbf{A} \mathbf{t} + \frac{1}{N} \mathbf{t}^T \mathbf{A} \mathbf{t} \quad (4)$$

$$= \frac{1}{N} \mathbf{w}^T \mathbf{X}^T \mathbf{A} \mathbf{X} \mathbf{w} - \frac{1}{N} \mathbf{t}^T \mathbf{A} \mathbf{X} \mathbf{w} - \frac{1}{N} \mathbf{w}^T \mathbf{X}^T \mathbf{A} \mathbf{t} + \frac{1}{N} \mathbf{t}^T \mathbf{A} \mathbf{t} \quad (5)$$

Since the two terms $\mathbf{t}^\top \mathbf{A} \mathbf{X} \mathbf{w}$ and $\mathbf{w}^\top \mathbf{X}^\top \mathbf{A} \mathbf{t}$ are the transpose of one another, and also scalars, we conclude that they must be the same and therefore we combine them.

$$\mathcal{L} = \frac{1}{N} \mathbf{w}^\top \mathbf{X}^\top \mathbf{A} \mathbf{X} \mathbf{w} - \frac{2}{N} \mathbf{w}^\top \mathbf{X}^\top \mathbf{A} \mathbf{t} + \frac{1}{N} \mathbf{t}^\top \mathbf{A} \mathbf{t} \quad (6)$$

$$= \frac{1}{N} (\mathbf{w}^\top \mathbf{X}^\top \mathbf{A} \mathbf{X} \mathbf{w} - 2 \mathbf{w}^\top \mathbf{X}^\top \mathbf{A} \mathbf{t} + \mathbf{t}^\top \mathbf{A} \mathbf{t}) \quad (7)$$

From this expression we find the derivative using the identities in Table 1.4.¹ Looking at the first term of \mathcal{L} we have $\mathbf{w}^\top \mathbf{X}^\top \mathbf{A} \mathbf{X} \mathbf{w}$. We know that $\mathbf{X}^\top \mathbf{X}$ gives a symmetric matrix. Since \mathbf{A} is a diagonal matrix we now that $\mathbf{X}^\top \mathbf{A} \mathbf{X}$ must still be symmetric. For this reason we can use the 4th identity from table 1.4 which gives us $2 \mathbf{X}^\top \mathbf{A} \mathbf{X} \mathbf{w}$. Looking at the second term of \mathcal{L} we have $2 \mathbf{w}^\top \mathbf{X}^\top \mathbf{A} \mathbf{t}$. We observe that $\mathbf{X}^\top \mathbf{A} \mathbf{t}$ equals a vector that does not depend on \mathbf{w} . Using the 1st identity from table 1.4 we get $\mathbf{X}^\top \mathbf{A} \mathbf{t}$. Looking at the last term of \mathcal{L} we see that $\mathbf{t}^\top \mathbf{A} \mathbf{t}$ does not depend on \mathbf{w} , which is why it disappears when deriving with respect to \mathbf{w} . From this we get the following derivative and solve for 0:

$$\frac{\partial \mathcal{L}}{\partial \mathbf{w}} = \frac{2}{N} \mathbf{X}^\top \mathbf{A} \mathbf{X} \mathbf{w} - \frac{2}{N} \mathbf{X}^\top \mathbf{A} \mathbf{t} = 0 \quad (8)$$

$$\mathbf{X}^\top \mathbf{A} \mathbf{X} \mathbf{w} = \mathbf{X}^\top \mathbf{A} \mathbf{t} \quad (9)$$

In order to isolate \mathbf{w} we will cancel $\mathbf{X}^\top \mathbf{A} \mathbf{X}$ out by premultiplying its inverse of both sides giving us:

$$\mathbf{I} \mathbf{w} = (\mathbf{X}^\top \mathbf{A} \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{A} \mathbf{t} \quad (10)$$

As $\mathbf{I} \mathbf{w} = \mathbf{w}$ we are left with a matrix equation for $\hat{\mathbf{w}}$, the value of \mathbf{w} that minimises the loss:

$$\hat{\mathbf{w}} = (\mathbf{X}^\top \mathbf{A} \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{A} \mathbf{t} \quad (11)$$

b)

Since we use $\alpha_n = t_n^2$ as the additional weights i expect that the big t_n values will have a bigger impact on the w_n values such that the spread will be less for high t_n values and and bigger for small t_n values. This is indeed what we observe when we look at figure 1.

See section 1.2 to see code for the solution.

¹S. Rogers Mark Girolami, A First Course In Machine Learning, 2nd edition, p.23




Figure 1: Figures containing scatter plot for true house prices vs estimated house prices. On the left is the plot from assignment 1, to the right is the plot using the weighted average loss found using the derivation found in a).

Exercise 2

a)

The best value of $\lambda = 0.0000003430$ and its $loss = 0.0624312040$

For this model $w_0 = 36.3776282284$ and $w_1 = -0.0133110046$

For the model with $\lambda = 0$, $w_0 = 36.4164559025$ and $w_1 = -0.0133308857$




Figure 2

b)

The best value of $\lambda = 0.0000205651$ and its $loss = 0.0521430656$

For this model $w_0 = 0.0188300350$, $w_1 = 9.1127782080$, $w_2 = -0.0138600662$, $w_3 = 0.0000070338$, $w_4 = -0.0000000012$

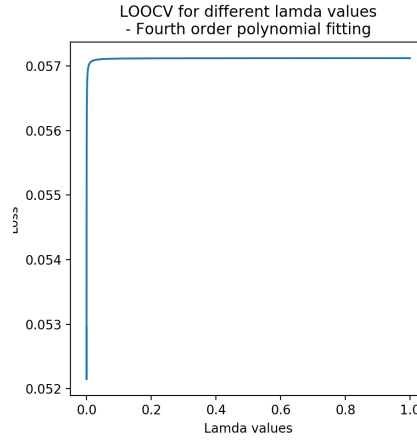


Figure 3

Exercise 3

a)

We are given the CDF:

$$F(x) = \begin{cases} 0 & \leq 0 \\ 1 - e^{-\beta x^\alpha} & > 0 \end{cases}$$

Since the CDF $F(X)$ is a continuous functions, we obtain the PDF by taking its derivative. Therefore we have

$$F'(x) = \begin{cases} 0 & \leq 0 \\ \beta \alpha e^{-\beta x^\alpha} x^{\alpha-1} & > 0 \end{cases}$$

b)

With $\alpha = 2$ and $\beta = \frac{1}{4}$. To find the probability that the chip works long than four years we use the CDF to compute

$$1 - F(4) \approx 0.0183156388873422$$

To find the probability that the chip stops working in the time interval $[5; 10]$ years we use the CDF to compute

$$F(10) - F(4) \approx 0.0019304541223398308$$

Exercise 4

a)

Given a person with no history of conviction (NC). The expected mean sentence duration he will have to spend in prison if

- NC talks to the police: $0.002 * (1 - 0.5) * 5 * 0.75 = 0.00375$ years
- NC stays silent: $0.001 * (1 - (0.5 * \frac{1}{4})) * 5 = 0.004375$ years

Given a person with a history of conviction (C). The expected mean sentence duration he will have to spend in prison is:

- C talks to the police: $0.005 * (1 - 0.1) * 0.5 * 0.75 = 0.003125$
- C stays silent: $0.001 * (1 - (0.1 * \frac{1}{4})) * 5 = 0.004875$

1 Appendix

1.1 Exercise 1 files

```

1000 import numpy as np
1002 class LinearWeightedRegression():
1003     def __init__(self):
1004
1005         pass
1006
1007     def fit(self, X, t):
1008         # Make sure that we have N-dimensional Numpy arrays (
1009         ndarray)
1010         X = np.array(X).reshape((len(X), -1))
1011         t = np.array(t).reshape((len(t), 1))
1012
1013         # Prepend a column of ones til the X matrix
1014         ones = np.ones((X.shape[0], 1))
1015         X = np.concatenate((ones, X), axis=1)
1016
1017         # Transform the w vector into a diagonal vector
1018         A = np.diag(t[:,0] ** 2)
1019
1020         self.w = np.linalg.solve(X.T @ A @ X, X.T @ A @ t)
1021
1022     def predict(self, X):
1023         X = np.array(X).reshape((len(X), -1))
1024
1025         ones = np.ones((X.shape[0], 1))
1026         X = np.concatenate((ones, X), axis=1)

```

```

1026         predictions = np.dot(X, self.w)
1028         return predictions

```

lineweighreg.py

```

1000 import numpy as np
1001 import lineweighreg
1002 import matplotlib.pyplot as plt
1003
1004 # Load data
1005 train_data = np.loadtxt("boston_train.csv", delimiter=",")
1006 test_data = np.loadtxt("boston_test.csv", delimiter=",")
1007
1008 # Choose all rows and all cols minus the last, choose all rows and
1009 # the last col
1010 X_train, t_train = train_data[:, :-1], train_data[:, -1]
1011 X_test, t_test = test_data[:, :-1], test_data[:, -1]
1012
1013 # Make sure that we have N-dimensional Numpy arrays (ndarray)
1014 t_train = t_train.reshape((len(t_train), 1))
1015 t_test = t_test.reshape((len(t_test), 1))
1016
1017 # Fit the model
1018 model = lineweighreg.LinearWeightedRegression()
1019 model.fit(X_train, t_train)
1020
1021 # Make predictions
1022 predictions = model.predict(X_test)
1023
1024 # Plot the data
1025 plt.figure(figsize=(5,5))
1026 plt.scatter(t_test, predictions, s=3)
1027 plt.title("A2")
1028 plt.xlabel("True house price")
1029 plt.ylabel("Estimate")
1030 plt.show()

```

exercisel.py

1.2 Exercise 2 files

```

1000 import numpy as np
1001
1002 class LinearRegression():
1003     def __init__(self, lam=0.0):
1004
1005         self.lam = lam
1006
1007     def fit(self, X, t):
1008
1009         X = np.array(X).reshape((len(X), -1))
1010         t = np.array(t).reshape((len(t), 1))

```

```

1012     ones = np.ones((X.shape[0], 1))
1013     X = np.concatenate((ones, X), axis=1)
1014
1015     diag = self.lam * len(X) * np.identity(X.shape[1])
1016     a = np.dot(X.T, X) + diag
1017     b = np.dot(X.T, t)
1018     self.w = np.linalg.solve(a, b)
1019
1020     def predict(self, X):
1021         X = np.array(X).reshape((len(X), -1))
1022
1023         ones = np.ones((X.shape[0], 1))
1024         X = np.concatenate((ones, X), axis=1)
1025
1026         predictions = np.dot(X, self.w)
1027
1028         return predictions

```

linereg.py

```

1000 import numpy as np
1001 import linereg
1002 import matplotlib.pyplot as plt
1003
1004 # DATA
1005 data = np.loadtxt("men-olympics-100.txt")
1006
1007 X, t = data[:, 0], data[:, 1]
1008 X = X.reshape((len(X), 1))
1009 t = t.reshape((len(t), 1))
1010
1011 lamdaValues = np.logspace(-8, 0, 100, base=10)
1012
1013 def LOOCV(X, t, l):
1014     errors = []
1015     for lam in lamdaValues:
1016         model = linereg.LinearRegression(lam=lam)
1017         modelError = 0
1018         for i in range(X.shape[0]):
1019             X_train = np.delete(X, i, 0)
1020             t_train = np.delete(t, i, 0)
1021
1022             X_pred = X[i]
1023             X_pred = X_pred.reshape((-1, len(X_pred)))
1024             t_pred = t[i]
1025             t_pred = t_pred.reshape((len(t_pred), 1))
1026
1027             model.fit(X_train, t_train)
1028
1029             prediction = model.predict(X_pred)
1030
1031             error = (prediction - t_pred)**2
1032             modelError += error[0][0]
1033         modelError = modelError / X.shape[0]
1034         errors.append(modelError)
1035     return errors
1036

```

```

# First order polynomial
1038 firstOrderLOOCV = LOOCV(X, t, lamdaValues)
    bestValueindex = np.argmin(firstOrderLOOCV)
1040 bestLamda = lamdaValues[bestValueindex]

1042 model = linereg.LinearRegression(lam=bestLamda)
    model.fit(X, t)
1044 print("===== First order polynomial fitting =====")
    print("The best value of lamda=%.10f and its loss=%.10f" %
1046         (bestLamda, firstOrderLOOCV[bestValueindex]))
    print("For this model w0=%.10f and w1=%.10f \n" %
1048         (model.w[0][0], model.w[1][0]))

1050 modelLam0 = linereg.LinearRegression()
    modelLam0.fit(X, t)
1052 print("For the model with lamda=0 w0=%.10f and w1=%.10f\n" %
        (modelLam0.w[0][0], modelLam0.w[1][0]))
1054

# Plot data
1056 plt.figure(figsize=(5,5))
    plt.plot(lamdaValues, firstOrderLOOCV)
1058 plt.title("LOOCV for different lamda values\n- First order
        polynomial fitting")
    plt.xlabel("Lamda values")
1060 plt.ylabel("Loss")
    plt.show()
1062

# Fourth order polynomial
1064 def augment(X, max_order):
    X_augmented = X

1066     for i in range(2, max_order+1):
1068         X_augmented = np.concatenate([X_augmented, X**i], axis=1)
    return X_augmented
1070

1072 Xnew = augment(X, 4)
    fourthOrderLOOCV = LOOCV(Xnew, t, lamdaValues)
1074 bestValueindex = np.argmin(fourthOrderLOOCV)
    bestLamda = lamdaValues[bestValueindex]
1076

    modelFourthOrder = linereg.LinearRegression(lam=bestLamda)
    modelFourthOrder.fit(Xnew, t)
1078 print("===== Fourth order polynomial fitting =====")
    print("The best value of lamda=%.10f and its loss=%.10f" %
1080         (bestLamda, fourthOrderLOOCV[bestValueindex]))
1082 print("For this model w0=%.10f, w1=%.10f, w2=%.10f, w3=%.10f, w4
        =%.10f \n" %
1084         (modelFourthOrder.w[0][0], modelFourthOrder.w[1][0],
            modelFourthOrder.w[2][0], modelFourthOrder.w[3][0],
            modelFourthOrder.w[4][0]))
1086

# Plot data
1088 plt.figure(figsize=(5,5))
    plt.plot(lamdaValues, fourthOrderLOOCV)
1090 plt.title("LOOCV for different lamda values\n- Fourth order
        polynomial fitting")

```



```
1092 plt.xlabel("Lamda values")  
plt.ylabel("Loss")  
plt.show()
```

exercise2.py